


DEPARTMENT OF INFORMATICS


TECHNISCHE UNIVERSITÄT MÜNCHEN

Bachelor's Thesis in Informatics

Multiple Precision Floating Point Arithmetic in Isabelle/HOL

Fabian Hellauer





DEPARTMENT OF INFORMATICS


TECHNISCHE UNIVERSITÄT MÜNCHEN

Bachelor's Thesis in Informatics

Multiple Precision Floating Point Arithmetic in
Isabelle/HOL

Multiple Precision Floating Point Arithmetic in
Isabelle/HOL

Author:	Fabian Hellauer
Supervisor:	Prof. Tobias Nipkow
Advisor:	Fabian Immler
Submission date:	March 15, 2016



I confirm that this bachelor's thesis is my own work and I have documented all sources and material used.

München, March 15, 2016

Fabian Hellauer

Abstract

Many problems in geometry or topology systems can only be approximated by large computations with long sequences of operations. To best benefit from the speedup of machine computation, one has to make use of the hardware floating point instructions that are specified by the "Institute of Electrical and Electronics Engineers" (IEEE) 754 standard[1]. However, using the IEEE-floats directly would often create the need for a complicated numerical analysis due to them being affected by round-off. Another one of their properties removes this need however: Alongside the result, the round-off error of an addition or subtraction can be computed. Storing and using it in further operations thus makes the computation error-free. We give a simple data format in Isabelle/HOL [8] that uses this approach to provide fast algorithms for error-free addition, subtraction and multiplication.

Acknowledgments

Fabian Immler provided a highly professional guidance and an incredible amount of support. His versatility for providing solutions in different aspects of scientific work inspired me to invest much time and energy aiming for a deeper understanding of automated theorem proving.

Tobias Nipkow's lectures taught me concepts of functional programming and motivated me to take the course in Semantics and apply for a bachelor's thesis to his research group.

Contents

Abstract	vii
Acknowledgments	ix
1 Introduction	1
1.1 Round-off	1
1.2 Problem statement	2
1.3 Contributions	2
2 Code Analysis	3
2.1 IEEE in Isabelle	3
2.2 Notation	3
2.3 Making operations error-free	4
2.3.1 Addition	4
2.3.2 Subtraction	5
2.4 Code Analysis	5
2.4.1 Tail recursive version	11
2.5 Further Operations	13
3 Code generation	15
3.1 Use of SML floats	15
3.2 Printing Floats	15
4 Results and conclusion	19
4.1 Impact	19
4.2 Future Work	19
Bibliography	21

1 Introduction

When attacking computational problems by machine, often times fixed precision numbers have to be used: Arbitrary precision numbers are very slow, mostly because they don't use the hardware floating point operations that are available to modern systems. If the finite set of these machine numbers suffice (or an approximation within their range), the use of hard-wired operations can speed up the computation by a large factor. However, they introduce the problem of round-off, which, when not handled, will affect the output's precision in complex ways.

1.1 Round-off

Round-off occurs when the result of a floating-point operation cannot be represented as a datum of the same format. This happens in particular if the size of exponent and mantissa is fixed, as it is case for floats defined by the IEEE standard for floating-point arithmetic[1]. Round-off also occurs when casting measured or infinitely precise data to such a limited precision format.

Dealing with round-off If round-off affected arithmetic is used in a long sequence of operations, the result will only approximate within a certain range. Correctness proofs for assertions to this range will require a tedious numerical analysis of the algorithms which is very complex to do formally.

Avoiding round-off Another approach is to avoid round-off altogether. However, using an implementation of infinitely precise rationals might severely slow down the code's execution due to them not making use of hardware floating point operations that modern machines provide.

Float expansions If a finite set of numbers with magnitude and precision in the range of IEEE floats suffices however, both the precision and fast execution speed can be preserved: This thesis presents the "float expansion" approach, where the accumulated errors are stored in a list alongside with an approximation for the result of the executed sequence. It provides addition, subtraction and multiplication within the numbers representable in this way (a finite superset to IEEE floats).

1.2 Problem statement

Isabelle[8] already provides the arbitrary precision format *real* in its *Complex-Main* library. The widely popular "IEEE-floats"[1] are modelled in an Isabelle theory *IEEE-Floating-Point/IEEE*[12] provided by the "Archive of Formal Proofs" (AFP).

The task for this bachelor thesis is to use this formalization to present a "multiple precision" float arithmetic in Isabelle/HOL. In several scientific papers this is described as "floating point expansion"[9][5] or "multiple term"[5] strategy. It is an easy way to gain considerable amounts of precision while still using the IEEE floating point specification to enable the widely available acceleration of hardware operations.

1.3 Contributions

We explain different aspects of the "floating point expansion" approach and then provide the data format *mpf*, which stands for "multiple precision float". It implements error-free addition, subtraction and multiplication within the numbers representable in this format (a finite superset to IEEE floats). We use the formal setting of Isabelle/HOL to specify and prove the algorithms correct, but we make sure all of them can easily be executed by adapting Isabelle's Standard ML (SML) code generation for IEEE-floats.

2 Code Analysis

2.1 IEEE in Isabelle

The IEEE standard for floating point arithmetic (IEEE 754-2008)[1] is already modelled in Lei Yu’s AFP entry *IEEE-Floating-Point/IEEE*. The formalization is quite general to accommodate for the many different allowed formats that arise when different precisions and exponent ranges are combined (the decimal formats are omitted). However, a strong precedence for the ”binary64” format and the ”roundTiesToEven” rounding mode can be observed. The operations using this format and rounding rule (called *float-format* respectively *To-nearest* in the theory) are wrapped into definitions with simpler names, e.g.

definition *plus-float* :: *float* \Rightarrow *float* \Rightarrow *float* **where** $a + b = \text{Abs-float } (fadd \text{ float-format } To\text{-nearest } (Rep\text{-float } a) (Rep\text{-float } b))$

In the case of the format, this is justifiable by ”binary64” being widely popular and hardware-implemented on most systems.

In the case of the rounding mode, the IEEE standard defines ”roundTiesToEven” to be the default ([1]p. 16).

We also use this format and rounding mode, which enables us to use the code printing defined in the theory *Code-Float* from the same AFP entry. The ”roundTiesToEven” rule is explicitly required for the TwoSum-properties ([7] sect. 4.3.3) and thus for all the presented algorithms.

2.2 Notation

In theory *IEEE*, the float operations use the $+$ and $-$ sign via

instantiation *float*

For this thesis, we want the IEEE-operations to use a different symbol, and reserve the $+$ operator for exact number formats. Fortunately, Isabelle provides spare symbols and the **abbreviation** command, which makes sure even the output uses the new notation. We thus state at the beginning of our **theory**:

— Use another notation for the possibly inexact IEEE-operations.

abbreviation *round-affected-plus* :: *float* \Rightarrow *float* \Rightarrow *float* (**infixl** \oplus 65) **where**
round-affected-plus $a\ b \equiv a + b$

abbreviation *round-affected-minus* :: *float* \Rightarrow *float* \Rightarrow *float* (**infixl** \ominus 65) **where**
round-affected-minus $a\ b \equiv a - b$

2 Code Analysis

Afterwards,

term $a + (b::float)$

outputs

$a \oplus b :: IEEE.float$

2.3 Making operations error-free

The core idea is to provide an error-free form of the basic operations between IEEE floats. Since we want the output to be floats as well, and rounding occurs for almost all input values, the only way to do so is to use the round-affected IEEE operation and then computing the error, also represented as floats. For the basic operations, this will turn out to be exactly another float.

2.3.1 Addition

We compute $a \oplus b$ together with y the error value. y will have the sign $-$ or $+$ corresponding to whether $a \oplus b$ is above resp. below the exact mathematical result of $a + b$. It can be computed by the following sequence, first described by Ole Møller in 1965[6]:

definition $TwoSum :: float \Rightarrow float \Rightarrow float \times float$ **where**

```
TwoSum a b = (let
  x = a  $\oplus$  b;
  bv = x  $\ominus$  a;
  av = x  $\ominus$  bv;
  br = b  $\ominus$  bv;
  ar = a  $\ominus$  av;
  y = ar  $\oplus$  br
in (x, y))
```

Here, we compute a value y such that $a + b = x + y$, where $x = a \oplus b$. The following lemma states the latter:

lemma *TwoSum-correct1*: $TwoSum\ a\ b = (x, y) \implies x = a \oplus b$

— x is defined in the first line of *TwoSum* and not changed thereafter.

by (*auto simp: TwoSum-def Let-def*)

The other property needs the preconditions that both the input and the output represent real numbers (as opposed to the special values *NaN* and $\pm\infty$). This is checked by the predicate *Finite*. We use the exact arithmetic of *real* and the conversion $Val :: float \Rightarrow real$ from *IEEE-Floating-Point/IEEE*.

lemma *TwoSum-correct2*:

```
fixes a b x y :: float
assumes Finite a
assumes Finite b
```



```

assumes Finite ( $a \oplus b$ )
assumes out:  $(x, y) = \text{TwoSum } a \ b$ 
shows  $\text{Val } a + \text{Val } b = \text{Val } x + \text{Val } y$ 
sorry

```

We assume the lemma by **sorry** for this thesis. Notice that a formal proof using the theorem prover Coq[2] is available online[11]

2.3.2 Subtraction

For $a - b$, we can compute:

definition *TwoDiff* :: *float* \Rightarrow *float* \Rightarrow *float* \times *float* **where**
TwoDiff $a \ b = \text{TwoSum } a \ (\text{float-neg } b)$

To drop the additional negation step, we could instead perform:

```

TwoDiff'  $a \ b = (\text{let}$ 
   $x = a \ominus b;$ 
   $b_v = x \ominus a;$ 
   $a_v = x \oplus b_v;$ 
   $b_r = b_v \ominus b;$ 
   $a_r = a \ominus a_v;$ 
   $y = a_r \oplus b_r$ 
   $\text{in } (x, y))$ 

```

according to Shewchuk[10]. Note that we still have a $+$ on the right side of the equation:

lemma *TwoDiff-correct2*:
fixes $a \ b \ x \ y :: \text{float}$
assumes *Finite* a
assumes *Finite* b
assumes *Finite* ($a \ominus b$)
assumes *out*: $(x, y) = \text{TwoDiff } a \ b$
shows $\text{Val } a - \text{Val } b = \text{Val } x + \text{Val } y$
sorry

Furthermore, Dekker[3] shows that in some situations, a sequence of three operations suffices. In order to not further increase the amount of unproven lemmas (or complicate them), we drop this optimization and the *TwoDiff* sequence for this thesis. This keeps the possibilities for errors at a minimum.

2.4 Code Analysis

The new data format is designed to implement the idea of storing all the errors as an unevaluated sum. It is defined as follows:

— Define the "Multiple Precision Float"
type-synonym *mpf* = *float* \times *float list*

2 Code Analysis

```
fun approx :: mpf  $\Rightarrow$  float where
  approx (a, es) = a
fun errors :: mpf  $\Rightarrow$  float list where
  errors (a, es) = es
```

where the tuple of a *float* and a *float list* should be seen together as a non-empty float list, ordered by decreasing magnitude. The approximation *approx* is just stored separately to avoid having to check for an empty list on access. Its quality depends on the executed algorithms (proofs about it can be found at [10]).

The mpf's represented value is the infinite precise sum of all its components:

```
fun Val-mpf :: mpf  $\Rightarrow$  real where
  Val-mpf (a, es) = Val a + listsum (map Val es)
```

Note that multiple *mpfs* can represent the same value. Many of these are invalid if we enforce the "non-overlapping" property proposed by Shewchuk[10] on the float list. However, since it needs to read out the bit representation of the IEEE float, there is no easy way to check this condition using the AFP-formalization. We could instead decrease the number of valid representations by not allowing zero components in the list's tail:

```
fun valid :: mpf  $\Rightarrow$  bool where
  valid (a, es) = (case Iszero a of
    True  $\Rightarrow$  es = [] |
    False  $\Rightarrow$  Finite a  $\wedge$  list-all ( $\lambda f$ . Isdenormal f  $\vee$  Isnormal f) es)
```

where λf . Isdenormal f \vee Isnormal f returns *False* for zero-floats:

```
lemma Iszero fl  $\implies \neg(\lambda f$ . Isdenormal f  $\vee$  Isnormal f) fl
using float-distinct
by (metis Isnormal-def Iszero-def is-normal-def is-zero-def order-less-irrefl)
```

Since zero components don't contribute to the mpf's value, omitting them is an easy way to save storage by decreasing the list size. The problem with this property is that the algorithms don't preserve the constraint by default. As Shewchuk puts it:

"A complicating characteristic of all the algorithms for manipulating expansions is that there may be spurious zero components scattered throughout the output expansions, even if no zeros were present in the input expansions."

As he shows by an example, they even occur in the middle of output lists that provably have all non-zero-components sorted. He also states:

"Unfortunately, accounting for these zero components could complicate the correctness proofs significantly." [10]

In other words: We **could** modify the algorithms to drop the zero component on-the-fly, but the extra branch would drastically increase the proof size. We instead settle for an even weaker property:

```
fun Finite-mpf :: mpf  $\Rightarrow$  bool where
  Finite-mpf (a, es)  $\longleftrightarrow$  Finite a  $\wedge$  list-all Finite es
```

Using this property and the well-known *TwoSum-correct2* property described above ($\llbracket \text{Finite } ?a; \text{Finite } ?b; \text{Finite } (?a \oplus ?b); (?x, ?y) = \text{TwoSum } ?a ?b \rrbracket \implies \text{Val } ?a + \text{Val } ?b = \text{Val } ?x + \text{Val } ?y$, unproven in this thesis), we will be able to prove computations error-free in the next section.

Here is a proof that valid implies Finite:

```
lemma valid-finite: valid (a, es)  $\implies$  Finite-mpf (a, es)
apply (simp split: bool.splits)
using float-cases-finite float-distinct apply fastforce
by (metis (no-types, lifting) Ball-set Finite-def)
```

```
definition safe-TwoSum a b =
  (let r = TwoSum a b in
   if Finite (fst r)  $\wedge$  Finite (snd r)
   then Some r
   else None)
```

```
definition safe-TwoDiff a b =
  (let r = TwoDiff a b in
   if Finite (fst r)  $\wedge$  Finite (snd r)
   then Some r
   else None)
```

```
lemma safe-TwoSum-finite:
assumes safe-TwoSum a b = Some (s, e)
shows safe-TwoSum-finite1: Finite s
and safe-TwoSum-finite2: Finite e
using assms
by (auto simp: safe-TwoSum-def Let-def split: split-if-asm)
```

```
lemma safe-TwoSum-correct1:
  safe-TwoSum a b = Some (x, y)  $\implies$  x = a  $\oplus$  b
by (auto simp: safe-TwoSum-def Let-def TwoSum-correct1 split: split-if-asm)
```

```
lemma safe-TwoSum-correct2:
fixes a b x y :: float
assumes Finite a Finite b Finite (a  $\oplus$  b)
assumes out: safe-TwoSum a b = Some (x, y)
shows Val a + Val b = Val x + Val y
using assms
by (auto intro!: TwoSum-correct2 simp: safe-TwoSum-def Let-def split: split-if-asm)
```

2 Code Analysis

definition $IsZero\text{-}mpf\ mpf \longleftrightarrow Iszero\ (approx\ mpf) \wedge errors\ mpf = []$

lemma $float\text{-}distinct\text{-}10: \neg (Isnormal\ f \wedge Iszero\ f)$
by $(auto\ simp\ add: float\text{-}defs\ is\text{-}normal\text{-}def\ is\text{-}zero\text{-}def)$

lemma $valid\text{-}no\text{-}zero\text{-}components: valid\ (a, es) \implies list\text{-}all\ (\lambda f. \neg Iszero\ f)\ es$
apply $(simp\ split: bool.splits)$
apply $(induction\ es)$
using $float\text{-}distinct(9)\ float\text{-}distinct\text{-}10$
apply $auto$
done

lemma $rec\text{-}val: Val\text{-}mpf\ (a, e \# es) = Val\ a + Val\text{-}mpf\ (e, es)$
by $simp$

lemma $rec\text{-}finite: Finite\text{-}mpf\ (a, e \# es) \longleftrightarrow Finite\ a \wedge Finite\text{-}mpf\ (e, es)$
by $simp$

fun $safe\text{-}grow\text{-}mpf\text{-}rec :: mpf \Rightarrow float \Rightarrow mpf\ option$ **where**
 $safe\text{-}grow\text{-}mpf\text{-}rec\ (a, [])\ f =$
 $\quad do\ \{$
 $\quad\quad (x, y) \leftarrow safe\text{-}TwoSum\ f\ a;$
 $\quad\quad Some\ (x, [y])$
 $\quad\ \}$ |
 $safe\text{-}grow\text{-}mpf\text{-}rec\ (a, e \# es)\ f =$
 $\quad do\ \{$
 $\quad\quad (a', es') \leftarrow safe\text{-}grow\text{-}mpf\text{-}rec\ (e, es)\ f;$
 $\quad\quad (x, y) \leftarrow safe\text{-}TwoSum\ a'\ a;$
 $\quad\quad Some\ (x, y \# es')$
 $\quad\ \}$

At this point, we could implement the zero removal explained before, by modifying the last lines of the blocks:

fun $safe\text{-}grow\text{-}mpf\text{-}rec\text{-}no\text{-}0 :: mpf \Rightarrow float \Rightarrow mpf\ option$ **where**
 $safe\text{-}grow\text{-}mpf\text{-}rec\text{-}no\text{-}0\ (a, [])\ f =$
 $\quad do\ \{$
 $\quad\quad (x, y) \leftarrow safe\text{-}TwoSum\ f\ a;$
 $\quad\quad if\ Iszero\ y\ then\ Some\ (x, [])\ else\ Some\ (x, [y])$
 $\quad\ \}$ |
 $safe\text{-}grow\text{-}mpf\text{-}rec\text{-}no\text{-}0\ (a, e \# es)\ f =$
 $\quad do\ \{$
 $\quad\quad (a', es') \leftarrow safe\text{-}grow\text{-}mpf\text{-}rec\text{-}no\text{-}0\ (e, es)\ f;$
 $\quad\quad (x, y) \leftarrow safe\text{-}TwoSum\ a'\ a;$
 $\quad\quad if\ Iszero\ y\ then\ Some\ (x, es')\ else\ Some\ (x, y \# es')$
 $\quad\ \}$

However, we don't pursue this idea further due to the problems mentioned there.

We rename the induction cases:

lemmas *safe-grow-mpf-induct* = *safe-grow-mpf-rec.induct*[*case-names no-error in-between*]

lemma *preserve-finite*:

assumes *safe-grow-mpf-rec mpf x = Some r*

assumes *Finite x Finite-mpf mpf*

shows *Finite-mpf r*

using *assms*

proof (*induction mpf x arbitrary: r rule: safe-grow-mpf-induct*)

— The base case is the case where the mpf is a single float with an empty error-list:

case (*no-error a f*)

— We apply the definition of *safe-grow-mpf-rec*:

from *no-error.premis(1)* **have** *do {(x, y) ← safe-TwoSum f a; Some (x, [y])} = Some r*

unfolding *safe-grow-mpf-rec.simps(1)* .

— Since we required the result to be some value, we can give it a name:

then obtain *x y* **where** *xy: safe-TwoSum f a = Some (x, y)* **and** *r: r = (x, [y])*

by (*auto simp: bind-eq-Some-conv*)

— and then delegate to the corresponding property of *safe-TwoSum*:

moreover from *safe-TwoSum-finite[OF xy]*

have *Finite x Finite y*.

ultimately show *?case*

by *simp*

next

case (*in-between a e es f r-full*)

note *in-between.premis(1)[simplified, unfolded bind-eq-Some-conv, simplified]*

then obtain *l r* **where** *goal1: safe-grow-mpf-rec (e, es) f = Some (l, r)*

and *r1: do {(x, y) ← safe-TwoSum l a; Some (x, y # r)} = Some r-full*

by *blast*

then obtain *l2 r2* **where** *l2: safe-TwoSum l a = Some (l2, r2)* **and**

r2: (l2, r2 # r) = r-full

using *r1[unfolded bind-eq-Some-conv, simplified]* **by** *auto*

from *r2* **have** *?case = Finite-mpf (l2, r2 # r)* **by** *simp*

moreover have *Finite l2*

using *safe-TwoSum-finite1[OF l2]*.

moreover have *Finite r2*

using *safe-TwoSum-finite2[OF l2]*.

moreover from *in-between.IH[OF goal1 in-between.premis(2)]* **have** *list-all Finite r*

using *in-between.premis(3)* **by** *auto*

ultimately

show *?case*

by *simp*

qed

Notice that the "assignments" (\leftarrow) in a Monad like *mpf option* can also be written using the \gg -operator (*bind*) and λ -notation. This is also what Isabelle's state panel will output, e.g.

do {(x, y) ← safe-TwoSum l a; Some (x, y # r)}

becomes

2 Code Analysis

safe-TwoSum $l\ a \gg (\lambda(x, y). \text{Some } (x, y \# r))$

etc. We perform the next proof using this style:

lemma *preserve-val*:

assumes *safe-grow-mpf-rec* $\text{mpf } x = \text{Some } r$

assumes *Finite* x *Finite-mpf* mpf

shows $\text{Val-mpf } r = \text{Val-mpf } \text{mpf} + \text{Val } x$

using *assms*

proof (*induction* $\text{mpf } x$ *arbitrary*: r *rule*: *safe-grow-mpf-induct*)

case (*no-error* $a\ f$)

from *no-error.prem* $s(1)$ **have** *safe-TwoSum* $f\ a \gg (\lambda(x, y). \text{Some } (x, [y])) = \text{Some } r$

unfolding *safe-grow-mpf-rec.simps* (1) .

then obtain $x\ y$ **where** xy : *safe-TwoSum* $f\ a = \text{Some } (x, y)$ **and** r : $r = (x, [y])$

by (*auto simp: bind-eq-Some-conv*)

from *safe-TwoSum-finite1* $[OF\ xy]$

have *Finite* x .

from *no-error* **have** an : *Finite* a **by** *simp*

show *?case*

using *safe-TwoSum-correct2* $[OF\ \langle Finite\ f \rangle\ an - xy]\ \langle Finite\ x \rangle$

safe-TwoSum-correct1 $[OF\ xy]$

by (*auto simp: r split: prod.split*)

next

case (*in-between* $a\ e\ es\ f\ r\text{-full}$)

note *in-between.prem* $s(1)[simplified, unfolded\ bind-eq-Some-conv, simplified]$

then obtain $l\ r$ **where** *goal1*: *safe-grow-mpf-rec* $(e, es)\ f = \text{Some } (l, r)$

and $r1$: *safe-TwoSum* $l\ a \gg (\lambda(x, y). \text{Some } (x, y \# r)) = \text{Some } r\text{-full}$

by *blast*

then obtain $l2\ r2$ **where** $l2$: *safe-TwoSum* $l\ a = \text{Some } (l2, r2)$ **and**

$r2$: $(l2, r2 \# r) = r\text{-full}$

using $r1$ $[unfolded\ bind-eq-Some-conv, simplified]$ **by** *auto*

then have $\text{Val-mpf } r\text{-full} = \text{Val-mpf } (l2, r2 \# r)$ **by** *simp*

also have $\dots = \text{Val } l2 + \text{Val-mpf } (r2, r)$

by (*simp add: rec-val*)

also have $\dots = \text{Val } l2 + \text{Val } r2 + \text{listsum}(\text{map } \text{Val } r)$

by *simp*

also have $\dots = \text{Val } l + \text{Val } a + \text{listsum}(\text{map } \text{Val } r)$

proof –

from *in-between.prem* s **have** *Finite* l

using *goal1 preserve-finite* **by** *auto*

moreover have *Finite* a

using *in-between.prem* $s(3)$ **by** *simp*

moreover have *Finite* $(l + a)$

using $l2$ *safe-TwoSum-correct1 safe-TwoSum-finite1* **by** *auto*

moreover have $\text{Val } l + \text{Val } a = \text{Val } l2 + \text{Val } r2$

using *safe-TwoSum-correct2* $[OF\ \text{calculation } l2]$.

ultimately show *?thesis*

by *simp*

qed

finally show *?case*

```

using in-between goal1 rec-finite by auto
qed

```

Note that the proof tactics for *preserve-finite* and *preserve-val* are very similar (identical up to the **obtain** commands in both induction cases). They could be combined by stating

```

shows preserve-finite: Finite-mpf r
and preserve-val: Val-mpf r = Val-mpf mpf + Val x

```

in a single lemma with the same assumptions. However, to actually remove redundancy in the proofs, both goals would have to be combined again via

```

unfolding atomize-conj

```

Since the second result depends partly on the first one, many fact names (or alternatively: large HOL predicates combining unrelated facts) would accumulate during the proof. To maintain readability, we stick to the two-proof-solution. Instead of stating both goals in one lemma, we collect the results afterwards:

```

lemmas safe-grow-mpf-correct =
  preserve-finite
  preserve-val

```

2.4.1 Tail recursive version

We can also implement *grow-mpf* in a tail-recursive way. For simplicity, we drop the overflow-check via the *option* monad for now.

```

fun grow-mpf-it :: float list  $\Rightarrow$  float  $\Rightarrow$  float list  $\Rightarrow$  mpf where
  grow-mpf-it [] f hs = (f, hs) |
  grow-mpf-it (e # es) f hs = (let
    (x, y) = TwoSum f e
    in grow-mpf-it es x (y # hs))

```

This transformation was comparably easy because *grow-mpf* only needs one linear pass as the graphic shows.

```

fun grow-mpf-tr :: mpf  $\Rightarrow$  float  $\Rightarrow$  mpf where
  grow-mpf-tr (a, es) f = (let
    (a', es') = grow-mpf-it (rev es) f [];
    (x, y) = TwoSum a' a
    in (x, y # es'))

```

An interesting realisation is that the list recursion can instead be implemented using fold:

```

fun grow-mpf-step :: float  $\Rightarrow$  mpf  $\Rightarrow$  mpf where
  grow-mpf-step f (a, es) = (let
    (x, y) = TwoSum a f
    in (x, y # es))

```

2 Code Analysis

```
fun grow-by-fold :: mpf  $\Rightarrow$  float  $\Rightarrow$  mpf where
  grow-by-fold (a, es) f = foldr grow-mpf-step (a # es) (f, [])
```

We prepare an equivalence proof for *grow-mpf-tr* and *grow-by-fold* by providing some lemmas about *grow-mpf-tr* and *op @*.

```
lemma grow-it-append-accumulator:
  grow-mpf-it as f (hs @ hs') = (let
    (a, es) = grow-mpf-it as f hs
  in (a, es @ hs'))
apply (induction as arbitrary: f hs hs')
apply simp-all
apply (metis (no-types, lifting) Cons-eq-appendI case-prod-beta)
done
```

```
lemma grow-it-append-expansion:
  grow-mpf-it (as @ es) f hs = (let
    (a', es') = grow-mpf-it as f hs
  in grow-mpf-it es a' es')
apply (induction as arbitrary: f hs)
by (simp-all add: prod.case-eq-if)
```

Its effect on the wrapper *grow-mpf-tr* is

```
lemma grow-append-rev:
  grow-mpf-tr (a, es @ es') f = (let
    (a'', es'') = grow-mpf-it (rev es') f [];
    (a', es') = grow-mpf-it (rev es) a'' es'';
    (x, y) = TwoSum a' a
  in (x, y # es'))
by (simp add: case-prod-beta grow-it-append-expansion)
```

In case of an increase by a singleton, this can be simplified:

```
lemma grow-snoc-rev:
  (grow-mpf-tr (a, es @ [h]) f) = (let
    (x, y) = TwoSum f h;
    (a', es') = grow-mpf-it (rev es) x [y];
    (x', y') = TwoSum a' a
  in (x', y' # es'))
unfolding grow-append-rev[of a es [h] f]
apply simp
by (simp add: split-def)
```

The right part of the equation can also be written using *grow-mpf-tr*:

```
lemma gm-snoc1: (grow-mpf-tr (a, es @ [h]) f) = (let
  (x, y) = TwoSum f h;
  (a', es') = grow-mpf-tr (a, es) x
in (a', es' @ [y]))
by (induction es arbitrary: a) (simp-all add: case-prod-beta grow-it-append-expansion)
```


We expect compilers that optimize tail recursion to also optimize *foldr*. Thus, it is no longer necessary to make the functions tail recursive if they can be expressed by fold.

2.5 Further Operations

With

```
fun mpf-neg :: mpf  $\Rightarrow$  mpf where
  mpf-neg (a, es) = (float-neg a, map float-neg es)
```

lemma *valid-zero-mpf*:

shows *valid Plus-zero-mpf*

and *valid Minus-zero-mpf*

by (*simp-all add: Plus-zero-mpf-def Minus-zero-mpf-def float-zero1 float-zero2*)

One way to inspect which computations will be performed is to define a test mpf with dummy values and then use Isabelle's simplifier to apply the methods simplifications to the desired point:

definition $a_4 = \text{undefined}$

definition $a_3 = \text{undefined}$

definition $a_2 = \text{undefined}$

definition $a_1 = \text{undefined}$

definition $a_0 = \text{undefined}$

definition *test-mpf* = ($a_4, [a_3, a_2, a_1]$)

definition *output* = *grow-by-fold test-mpf a₀*

To make the simplifier apply the definition, we need to state a lemma:

lemma *P output unfolding output-def test-mpf-def grow-by-fold.simps*

— We can now use various proof methods to get a neatly arranged output:

apply (*clarsimp split: prod.splits*) **oops**

where P is an undefined dummy predicate. At the last step, the output is as follows:

$$\begin{aligned} &\bigwedge x1\ x2\ x1a\ x2a\ x1b\ x2b\ x1c\ x2c. \\ &\quad TwoSum\ x1b\ a_4 = (x1c, x2c) \implies \\ &\quad TwoSum\ x1a\ a_3 = (x1b, x2b) \implies \\ &\quad TwoSum\ x1\ a_2 = (x1a, x2a) \implies \\ &\quad TwoSum\ a_0\ a_1 = (x1, x2) \implies \\ &\quad P\ (x1c, [x2c, x2b, x2a, x2]) \end{aligned}$$

To demonstrate the *TwoSum* sequence carried out by *grow-mpf*, we use the following graphic:

value *approx output* delivers

Plus-zero \oplus *One* \oplus *undefined* \oplus *undefined* \oplus *undefined* \oplus *undefined* \oplus *undefined*
 $:: IEEE.float$

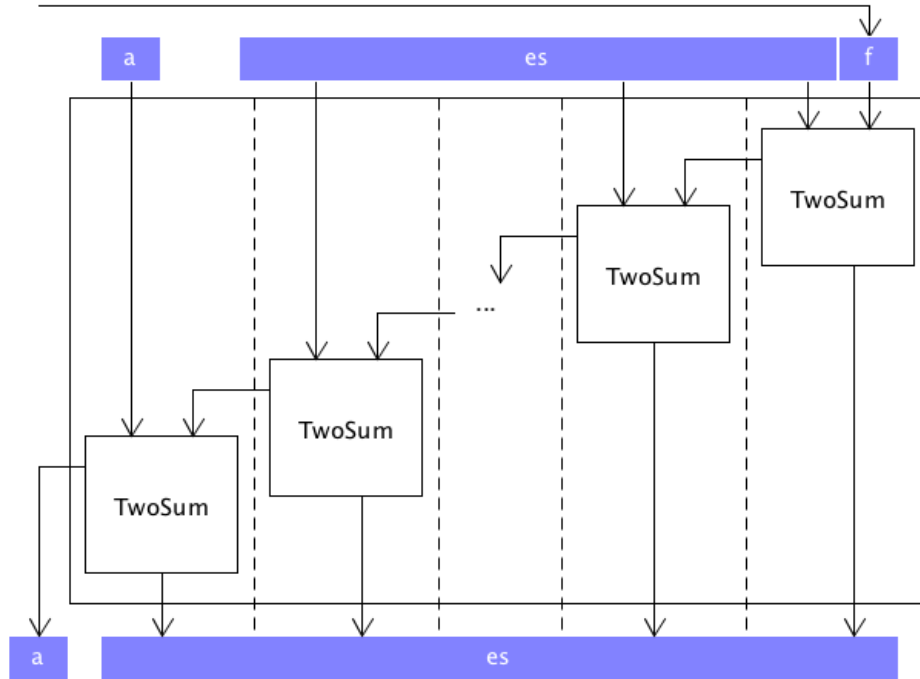


Figure 2.1: The float f is added to the mpf (a, es) . `TwoSum` is represented by a box where the larger value, x , is output to the left side and the smaller one, y , to the bottom. On the top, the function call is passed on. The returned mpf (bottom) is built from right to left.

3 Code generation

3.1 Use of SML floats

To enable computation for hardware floats, **theory** *Code-Float*[12] provides the built-in operators of the target language:

```
code-printing constant op / :: float  $\Rightarrow$  float  $\Rightarrow$  float  $\rightarrow$   
  (SML) Real.'/ ((-), (-)) and (OCaml) Pervasives.( '/' . )  
declare divide-float-def [code del]
```

The other operations are defined analogously.

Even ML's comparisons can be used (ML's `bool` is already defined as translation for HOL's *bool* in the *Main* theory *HOL*):

```
code-printing constant Orderings.less :: float  $\Rightarrow$  float  $\Rightarrow$  bool  $\rightarrow$   
  (SML) Real.< ((-), (-)) and (OCaml) Pervasives.( < )  
declare less-eq-float-def [code del]
```

3.2 Printing Floats

If we decide that an unchecked code module is safe enough for us, we can use the format *Float.float*[4] from Isabelle's HOL-library.

To enable the conversion from *IEEE.float* to *Float.float* in the generated code, we first insert the possibility to produce them from integers:

```
definition float-of-int i = Float (real-of-int i)  
context includes integer.lifting begin  
lift-definition float-of-integer::integer  $\Rightarrow$  float is float-of-int .  
end
```

```
lemma float-of-int[code]:  
  float-of-int i = float-of-integer (integer-of-int i)  
  by (simp add: float-of-integer-def)
```

```
code-printing  
  constant float-of-integer :: integer  $\Rightarrow$  float  $\rightarrow$  (SML) Real.fromInt  
declare [[code drop: float-of-integer]]
```

Then, the conversion is possible:

— convert hardware floats to *Float.float* for an exact representation

```
code-printing  
code-module ToManExp  $\rightarrow$  (SML)
```

3 Code generation

```

⟨fun tomanexp x =
  let
    val {man = m, exp = e} = Real.toManExp x;
    val p = Math.pow (2.0, 53.0);
    val ms = m * p;
    val mi = Real.floor ms;
    val ei = op Int.- (e, 53);
  in (mi, ei)
end⟩

```

consts *tomanexp::float* \Rightarrow *integer * integer*
code-printing constant *tomanexp :: float* \Rightarrow *integer * integer* \rightarrow
(SML) tomanexp

definition *toFloat::float* \Rightarrow *Float.float* **where**
toFloat x = (let (m, e) = tomanexp x in *Float.Float* (*int-of-integer* m) (*int-of-integer* e))

We can now define a test list:

definition *list :: float list* **where**
 — Note that floats with magnitude < 1 can only be defined via *op div*:
list = [
 float-of-int 43,
 float-of-int 34538,
 float-of-int 3 / *float-of-int* 44,
 float-of-int 0,
 float-of-int 0,
 float-of-int (−348976754389282980)]

To use the ML operators, we have to insert the transformation to a term and back:

instantiation *float::term-of*
begin
definition *term-of::float* \Rightarrow *term* **where** *term-of x* = *undefined*
instance ..
end

code-printing
code-module *FromManExp* \rightarrow (*SML*)
 ⟨fun frommanexp m e = *Real.fromManExp* {man = *Real.fromLargeInt* m, exp = e}⟩
consts *frommanexp::integer* \Rightarrow *integer* \Rightarrow *float*
code-printing constant *frommanexp :: integer* \Rightarrow *integer* \Rightarrow *float* \rightarrow
(SML) frommanexp

definition *of-Float::Float.float* \Rightarrow *float* **where**
of-Float x = *frommanexp* (*integer-of-int* (*Float.mantissa* x)) (*integer-of-int* (*Float.exponent* x))

lemma [*code*]: *term-of-class.term-of* (*x::float*) \equiv
Code-Evaluation.App

```

    (Code-Evaluation.termify of-Float)
    (term-of-class.term-of (normfloat (toFloat x)))
  by (rule term-of-anything)

```

We can now print the list without an error:

value *list*

produces

```

[of-Float 43, of-Float (Float.Float 17269 1),
 of-Float (Float.Float 1228254443828317 (- 54)),
 of-Float (Float.Float 0 0), of-Float (Float.Float 0 0),
 of-Float (Float.Float (- 5452761787332547) 6)]
:: float list

```

which is an error-free representation.

abbreviation *toNF* :: *float* ⇒ *Float.float* **where**
toNF ≡ *normfloat o toFloat*

3 Code generation

4 Results and conclusion

The IEEE 754 floats were already modelled in Isabelle. Using this formalization, this work provides an easy way to use them for fast and error-free addition and subtraction. For these operations, we translated algorithms from the literature and adapted them for our purposes.

4.1 Impact

We give a more practice-oriented analysis of Shewchuk’s algorithms and offer explanations for challenges that can arise when implementing them. We also give ideas and solutions for verifying them in a functional setting. Based on the existing formalization of IEEE-floats, we then specified a data format to provide an easy access to these algorithms. This means that users have a new option for a number format to perform verified computations using fast and error-free addition and subtraction. As results of our thorough testing of generated code, an error in polyML’s float handling has been detected and removed. This means the code generated using the AFP-theory *IEEE-Floating-Point/Code-Float* has now a clearer semantics.

4.2 Future Work

A correctness proof for the *TwoSum* method needs to be converted to Isabelle’s IEEE754 formalization. This will then also enable proofs for Shewchuk’s ”nonoverlapping” property, which, when implemented, allows more assertions about float expansions to be formally verified, e.g. about the maximum length of a valid *mpf*, or the quality of the approximation stored in the first component. Another improvement could be made by adapting code generation for IEEE-floats to support more of Isabelle’s target languages. This will make our arithmetic library more flexible for use in languages than SML. However, the correct behaviour of floats in the language should be ensured beforehand, to avoid getting wrong results when using the generated code.

Bibliography

- [1] IEEE Standard for Floating-Point Arithmetic. *IEEE Std 754-2008*, pages 1–70, Aug 2008. <http://ieeexplore.ieee.org/xpl/articleDetails.jsp?arnumber=4610935>.
- [2] Y. Bertot and P. Castéran. *Interactive Theorem Proving and Program Development. Coq'Art: The Calculus of Inductive Constructions*. Springer, 2004. <http://www.labri.fr/perso/casteran/CoqArt/index.html>.
- [3] T. J. Dekker. A floating-point technique for extending the available precision. *Numerische Mathematik*, 18(3):224–242.
- [4] J. Hölzl and F. Immler. Floating-point numbers. 2012. <https://isabelle.in.tum.de/library/HOL/HOL-Library/Float.html>, TU München.
- [5] M. Joldes, V. Popescu, and W. Tucker. Searching for sinks of Henon map using a multiple-precision GPU arithmetic library. Technical report, Nov. 2013. <https://hal.archives-ouvertes.fr/hal-00957438>, 7 pages.
- [6] O. Møller. Quasi double-precision in floating point addition. *BIT Numerical Mathematics*, 5(1):37–50.
- [7] J.-M. Muller, N. Brisebarre, F. de Dinechin, C.-P. Jeannerod, V. Lefèvre, G. Melquiond, N. Revol, D. Stehlé, and S. Torres. *Handbook of Floating-Point Arithmetic*. Birkhäuser Boston, 2010. ACM G.1.0; G.1.2; G.4; B.2.0; B.2.4; F.2.1., ISBN 978-0-8176-4704-9.
- [8] T. Nipkow, L. C. Paulson, and M. Wenzel. *Isabelle/HOL — A Proof Assistant for Higher-Order Logic*, volume 2283 of *LNCs*. Springer, 2002.
- [9] D. M. Priest. Algorithms for arbitrary precision floating point arithmetic. In *Computer Arithmetic, 1991. Proceedings., 10th IEEE Symposium on*, pages 132–143, Jun 1991.
- [10] J. R. Shewchuk. Adaptive Precision Floating-Point Arithmetic and Fast Robust Geometric Predicates. *Discrete & Computational Geometry*, 18(3):305–363, Oct. 1997.
- [11] L. Théry, L. Rideau, L. Fousse, G. Melquiond, and S. Boldo. Twosum. *A Coq Library on Floating-Point Arithmetic*. <http://lipforge.ens-lyon.fr/www/pff/TwoSum.html>.
- [12] L. Yu. A Formal Model of IEEE Floating Point Arithmetic. *Archive of Formal Proofs*, July 2013. http://afp.sf.net/entries/IEEE.Floating_Point.shtml, Formal proof development.